

Phase 2

• Discussion of the memory allocation capabilities using the following memory allocation policies:

1. First Fit
2. Next Fit
3. Best Fit
4. Buddy System Allocation

When a process is picked by the scheduler to be run, it must pass the memory check to get sure whether it can be allocated or it must wait till there is enough memory for it, the integration of this is in the next_ family functions, these functions now are responsible to get sure that there is some place in memory for a process if the scheduler wants to pick it, it sends that request to `allc(int mem)` function to ask for some amount of memory, `allc` in return replies with '-1' if there is no place in memory of that size and replies by the memory starting location if a place could be found, the `allc` function chooses the specified memory allocation algorithm and run it (FF,NF,BF,BSA).

Your algorithm explanation and results.

1. First Come First Serve (FCFS)

By imagining that process table as a circular array, the `pcb_ind` is the index of the last entered process in the process table and the `pcb_curr` is the index to the currently executing process, it is always the case that `pcb_curr` is behind or the same as `pcb_ind`, that assures that the processes will execute by the order of arrival.

2. Shortest Job First (SJF)

As this is a batch system, at first we start by picking the least running time process and after finishing we loop again and pick the least running time process and so on.

3. Preemptive Highest Priority First (HPF)

In this algorithm, we need to check for the highest process in terms of priority and execute it, whether there was a process executing before it or not, that is because this algorithm is preemptive so a process might be stopped to run another.

4. Shortest Remaining Time Next (SRTN)

Same as HPF but the condition over preemption is for the remaining time of the process which is kept track of by the process control block of this very process.

5. Round Robin (RR)

We use the quantum entered by the user to determine how much a process will run when it is picked by the scheduler, when no quantum time is entered by the user it will default to '1' clock cycle (1s).

The technique used is that when a process is running, any processes entering the system will reside at the end of the virtual queue data structure and when the quantum of the running process ends it will get back at the back of the queue.

TL;DR, when a process is stopped, it goes to the far back of the queue, back so that any incoming process at the same clock of the running process being stopped will reside before the stopped process in turn in the virtual queue.

***Note: the virtual queue mentioned is the same array described above, but it needs some functions to manipulate it to act like a queue for round robin to work as expected.**

• Your assumptions.

- 1- The maximum size of processes in the process table is 1024.
- 2- A process that takes zero time to execute instantaneously.
- 3- memory per process is at least 1 block.

• Workload distribution.

- 1- Mostafa Sobhy: Next Fit
- 2- Mostafa Magdy: First Fit
- 3- Zyad Hasan: Best Fit.
- 4- Omer Yacine: Buddy System Allocation

Test Cases

HPF (First Fit)

```
Makefile M [C] procs.m M x scheduler.log scheduler.perf memory.log
phase2 > code > exe > [C] procs.m
    You, seconds ago | 1 author (You)
1  #id arrival runtime priority memory
2  1 4 2 2 120
3  2 7 6 9 188
4  3 7 0 3 11
5  4 10 1 2 236
6  5 12 2 8 125
7

Makefile M [C] procs.m M x scheduler.log x scheduler.perf memory.log
phase2 > code > exe > scheduler.log
1  At time 4 process 1 started arr 4 total 2 remain 2 wait 0
2  At time 6 process 1 finished arr 4 total 2 remain 0 wait 0 TA 2 WTA 1.00
3  At time 7 process 3 started arr 7 total 0 remain 0 wait 0
4  At time 7 process 3 finished arr 7 total 0 remain 0 wait 0 TA 0 WTA 0.00
5  At time 7 process 2 started arr 7 total 6 remain 6 wait 0
6  At time 10 process 2 stopped arr 7 total 6 remain 3 wait 0
7  At time 10 process 4 started arr 10 total 1 remain 1 wait 0
8  At time 11 process 4 finished arr 10 total 1 remain 0 wait 0 TA 1 WTA 1.00
9  At time 11 process 2 resumed arr 7 total 6 remain 3 wait 1
10 At time 12 process 2 stopped arr 7 total 6 remain 2 wait 1
11 At time 12 process 5 started arr 12 total 2 remain 2 wait 0
12 At time 14 process 5 finished arr 12 total 2 remain 0 wait 0 TA 2 WTA 1.00
13 At time 14 process 2 resumed arr 7 total 6 remain 2 wait 3
14 At time 16 process 2 finished arr 7 total 6 remain 0 wait 3 TA 9 WTA 1.50
15

Makefile M [C] procs.m M x scheduler.log scheduler.perf x memory.log
phase2 > code > exe > scheduler.perf
1  CPU utilization = 81.25%
2  Avg WTA = 0.90
3  Avg Waiting = 0.60
4
```

```
Makefile M |C| procs.m M scheduler.log scheduler.perf memory.log x
phase2 > code > exe > memory.log
1   At time 4 allocated 120 bytes for process 1 from 0 to 119
2   At time 6 freed 120 bytes for process 1 from 0 to 119
3   At time 7 allocated 11 bytes for process 3 from 0 to 10
4   At time 7 freed 11 bytes for process 3 from 0 to 10
5   At time 7 allocated 188 bytes for process 2 from 0 to 187
6   At time 10 allocated 236 bytes for process 4 from 188 to 423
7   At time 11 freed 236 bytes for process 4 from 188 to 423
8   At time 12 allocated 125 bytes for process 5 from 188 to 312
9   At time 14 freed 125 bytes for process 5 from 188 to 312
10  At time 16 freed 188 bytes for process 2 from 0 to 187
11
```

Shortest Remaining Time Next (Next Fit)

```
Makefile M |C| procs.m M scheduler.log x scheduler.perf memory.log
phase2 > code > exe > scheduler.log
1   At time 4 process 2 started arr 4 total 0 remain 0 wait 0
2   At time 4 process 2 finished arr 4 total 0 remain 0 wait 0 TA 0 WTA 0.00
3   At time 4 process 3 started arr 4 total 0 remain 0 wait 0
4   At time 4 process 3 finished arr 4 total 0 remain 0 wait 0 TA 0 WTA 0.00
5   At time 4 process 1 started arr 4 total 1 remain 1 wait 0
6   At time 5 process 1 finished arr 4 total 1 remain 0 wait 0 TA 1 WTA 1.00
7   At time 7 process 4 started arr 7 total 1 remain 1 wait 0
8   At time 8 process 4 finished arr 7 total 1 remain 0 wait 0 TA 1 WTA 1.00
9   At time 9 process 5 started arr 9 total 4 remain 4 wait 0
10  At time 13 process 5 finished arr 9 total 4 remain 0 wait 0 TA 4 WTA 1.00
11
```

```
Makefile M [C] procs.m M x scheduler.log scheduler.perf memory.log
phase2 > code > exe > [C] procs.m
You, seconds ago | 1 author (You)
1 #id arrival runtime priority memory
2 1 4 1 10 129
3 2 4 0 1 197
4 3 4 0 2 181
5 4 7 1 10 82
6 5 9 4 2 9
7
```

As you can see, process 2 and 3 had runtime = 0 so they finished in the very same time clock that process 1 started, so actually process 1 didn't wait and started in its arriving clock cycle, this wait=0 & WTA=1.

```
Makefile M [C] procs.m M scheduler.log scheduler.perf x memory.log
phase2 > code > exe > scheduler.perf
1 CPU utilization = 69.23%
2 Avg WTA = 0.60
3 Avg Waiting = 0.00
4
```

```
Makefile M [C] procs.m M scheduler.log scheduler.perf memory.log x
phase2 > code > exe > memory.log
1 At time 4 allocated 197 bytes for process 2 from 129 to 325
2 At time 4 freed 197 bytes for process 2 from 129 to 325
3 At time 4 allocated 181 bytes for process 3 from 507 to 687
4 At time 4 freed 181 bytes for process 3 from 507 to 687
5 At time 4 allocated 129 bytes for process 1 from 817 to 945
6 At time 5 freed 129 bytes for process 1 from 817 to 945
7 At time 7 allocated 82 bytes for process 4 from 0 to 81
8 At time 8 freed 82 bytes for process 4 from 0 to 81
9 At time 9 allocated 9 bytes for process 5 from 82 to 90
10 At time 13 freed 9 bytes for process 5 from 82 to 90
11
```

RR (q=4)(Best Fit)

```
Makefile M  [C] procs.m M x scheduler.log scheduler.perf memory.log
phase2 > code > exe > [C] procs.m
You, seconds ago | 1 author (You)
1 #id arrival runtime priority memory
2 1 4 2 3 13
3 2 5 0 7 22
4 3 5 2 9 159
5 4 8 2 2 156
6 5 10 1 9 174
7 6 11 6 1 208
8 7 11 2 0 141
9

Makefile M  [C] procs.m M scheduler.log x scheduler.perf memory.log
phase2 > code > exe > scheduler.log
1 At time 4 process 1 started arr 4 total 2 remain 2 wait 0
2 At time 6 process 1 finished arr 4 total 2 remain 0 wait 0 TA 2 WTA 1.00
3 At time 6 process 2 started arr 5 total 0 remain 0 wait 1
4 At time 6 process 2 finished arr 5 total 0 remain 0 wait 1 TA 1 WTA 1.00
5 At time 6 process 3 started arr 5 total 2 remain 2 wait 1
6 At time 8 process 3 finished arr 5 total 2 remain 0 wait 1 TA 3 WTA 1.50
7 At time 8 process 4 started arr 8 total 2 remain 2 wait 0
8 At time 10 process 4 finished arr 8 total 2 remain 0 wait 0 TA 2 WTA 1.00
9 At time 10 process 5 started arr 10 total 1 remain 1 wait 0
10 At time 11 process 5 finished arr 10 total 1 remain 0 wait 0 TA 1 WTA 1.00
11 At time 11 process 6 started arr 11 total 6 remain 6 wait 0
12 At time 15 process 6 stopped arr 11 total 6 remain 2 wait 0
13 At time 15 process 7 started arr 11 total 2 remain 2 wait 4
14 At time 17 process 7 finished arr 11 total 2 remain 0 wait 4 TA 6 WTA 3.00
15 At time 17 process 6 resumed arr 11 total 6 remain 2 wait 2
16 At time 19 process 6 finished arr 11 total 6 remain 0 wait 2 TA 8 WTA 1.33
17

Makefile M  [C] procs.m M scheduler.log scheduler.perf x memory.log
phase2 > code > exe > scheduler.perf
• 1 CPU utilization = 84.21%
2 Avg WTA = 1.40
3 Avg Waiting = 1.14
4
```

```
Makefile M | C | procs.m M | scheduler.log | scheduler.perf | memory.log x
phase2 > code > exe > memory.log
1   At time 4 allocated 13 bytes for process 1 from 0 to 12
2   At time 6 freed 13 bytes for process 1 from 0 to 12
3   At time 6 allocated 22 bytes for process 2 from 0 to 21
4   At time 6 freed 22 bytes for process 2 from 0 to 21
5   At time 6 allocated 159 bytes for process 3 from 0 to 158
6   At time 8 freed 159 bytes for process 3 from 0 to 158
7   At time 8 allocated 156 bytes for process 4 from 0 to 155
8   At time 10 freed 156 bytes for process 4 from 0 to 155
9   At time 10 allocated 174 bytes for process 5 from 0 to 173
10  At time 11 freed 174 bytes for process 5 from 0 to 173
11  At time 11 allocated 208 bytes for process 6 from 0 to 207
12  At time 15 allocated 141 bytes for process 7 from 208 to 348
13  At time 17 freed 141 bytes for process 7 from 208 to 348
14  At time 19 freed 208 bytes for process 6 from 0 to 207
15
```

HPF (Buddy System Allocation)

```
Makefile M | C | procs.m M x | scheduler.log | scheduler.perf
phase2 > code > exe > C | procs.m
    You, seconds ago | 1 author (You)
1   #id arrival runtime priority memory
2   1 3 0 9 202
3   2 5 4 10 32
4   3 7 4 2 83
5   4 9 3 1 69
6   5 12 0 8 33
7
```

```
Makefile M  [C] procs.m M  scheduler.log x  scheduler.perf
phase2 > code > exe > scheduler.log
1  At time 3 process 1 started arr 3 total 0 remain 0 wait 0
2  At time 3 process 1 finished arr 3 total 0 remain 0 wait 0 TA 0 WTA 0.00
3  At time 5 process 2 started arr 5 total 4 remain 4 wait 0
4  At time 7 process 2 stopped arr 5 total 4 remain 2 wait 0
5  At time 7 process 3 started arr 7 total 4 remain 4 wait 0
6  At time 9 process 3 stopped arr 7 total 4 remain 2 wait 0
7  At time 9 process 4 started arr 9 total 3 remain 3 wait 0
8  At time 12 process 4 finished arr 9 total 3 remain 0 wait 0 TA 3 WTA 1.00
9  At time 12 process 3 resumed arr 7 total 4 remain 2 wait 3
10 At time 14 process 3 finished arr 7 total 4 remain 0 wait 3 TA 7 WTA 1.75
11 At time 14 process 5 started arr 12 total 0 remain 0 wait 2
12 At time 14 process 5 finished arr 12 total 0 remain 0 wait 2 TA 2 WTA 2.00
13 At time 14 process 2 resumed arr 5 total 4 remain 2 wait 7
14 At time 16 process 2 finished arr 5 total 4 remain 0 wait 7 TA 11 WTA 2.75
15
```

```
Makefile M  [C] procs.m M  scheduler.log  scheduler.perf x
phase2 > code > exe > scheduler.perf
1  CPU utilization = 81.25%
2  Avg WTA = 1.50
3  Avg Waiting = 2.40
4
```

```
Makefile M  [C] procs.m M  scheduler.log  scheduler.perf  memory.log x
phase2 > code > exe > memory.log
1  At time 3 allocated 256 bytes for process 1 from 0 to 255
2  At time 3 freed 256 bytes for process 1 from 0 to 255
3  At time 5 allocated 32 bytes for process 2 from 0 to 31
4  At time 7 allocated 128 bytes for process 3 from 128 to 255
5  At time 9 allocated 128 bytes for process 4 from 256 to 383
6  At time 12 freed 128 bytes for process 4 from 256 to 383
7  At time 14 freed 128 bytes for process 3 from 128 to 255
8  At time 14 allocated 64 bytes for process 5 from 64 to 127
9  At time 14 freed 64 bytes for process 5 from 64 to 127
10 At time 16 freed 32 bytes for process 2 from 0 to 31
11
```